

# Probabilistic Forecasting with `nnet`sauce (using Density Estimation, Bayesian inference, Conformal prediction and Vine copulas)

T. Moudiki ([thierrymoudiki.github.io](https://thierrymoudiki.github.io))

2024-07-26

# Context

- ▶ **Quasi-randomized *neural* networks** (QRNs) applied to time series lags for forecasting
- ▶ **Uncertainty quantification** using Kernel Density Estimation, Bayesian inference, Conformal prediction and Vine copulas
- ▶ Implemented in Python package **`nnetsauce` version 0.23.0**

# Plan

- ▶ 1 - Key components of `nnetsauce` forecasting
  - ▶ 1 - 1 Quasi-randomized *neural* networks (QRNs)
  - ▶ 1 - 2 Uncertainty quantification in forecasting
- ▶ 2 - QRN forecasting with `nnetsauce`
  - ▶ 2 - 1 `nnetsauce`'s description (Python version)
  - ▶ 2 - 2 Install+import Python packages (including `nnetsauce`)
  - ▶ 2 - 3 Import data for the demo
  - ▶ 2 - 4 Using the `fit + predict` interface
  - ▶ 2 - 5 Using GPUs
  - ▶ 2 - 6 Time series cross-validation
  - ▶ 2 - 7 AutoML with LazyMTS

# Plan

- ▶ **1 - Key components of `nnetsauce` forecasting**
  - ▶ 1 - 1 Quasi-randomized *neural* networks (QRNs)
  - ▶ 1 - 2 Uncertainty quantification in forecasting
- ▶ **2 - QRN forecasting with `nnetsauce`**
  - ▶ 2 - 1 `nnetsauce`'s description (Python version)
  - ▶ 2 - 2 Install+import Python packages (including `nnetsauce`)
  - ▶ 2 - 3 Import data for the demo
  - ▶ 2 - 4 Using the `fit + predict` interface
  - ▶ 2 - 5 Using GPUs
  - ▶ 2 - 6 Time series cross-validation
  - ▶ 2 - 7 AutoML with LazyMTS

# 1 - Key components of nnet sauce forecasting

## 1 - 1 Quasi-randomized *neural* networks (QRNs)

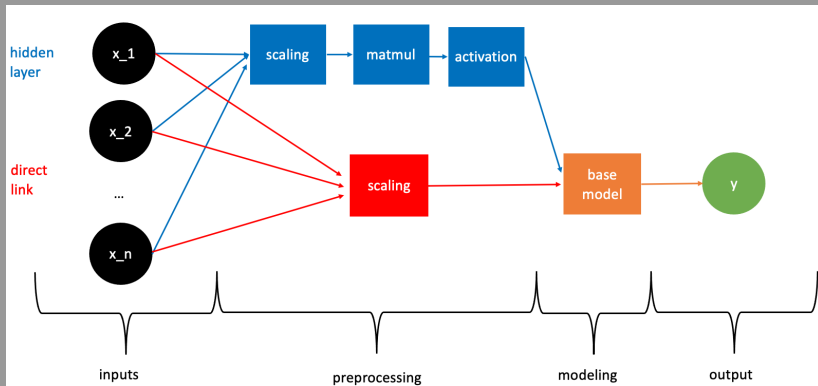


Figure 1: QRN principle

- ▶ See also [this doc](#)
- ▶ A simple case: **base learner** = Linear Regression (next page)

## 1 - 1 - Quasi-randomized *neural* networks (QRNs)

**Simple case:** **base learner** = Linear Regression;  $y \in \mathbb{R}^n$ , to be explained by  $X^{(j)}, j \in \{1, \dots, p\}$

$$y = \beta_0 + \sum_{j=1}^p \beta_j x^{(j)} + \sum_{l=1}^L \gamma_l g \left( \sum_{j=1}^p w^{(j,l)} x^{(j)} \right) + \epsilon$$

With:

- ▶  $g$  : **activation function**  $\rightarrow$  nonlinearity
- ▶  $L$  : number of nodes in the **hidden layer**
- ▶  $w^{(j,l)}$ , hidden layer: **pseudo/quasi-random**
- ▶ **Quasi-random**: designed to cover the space *parsimoniously*
- ▶ **secret sauce**: “*Layer normalization*” (centering and scaling twice)
- ▶  $\beta_j$  and  $\gamma_l$  : linear model coefficients
- ▶  $\epsilon$  : residuals

## 1 - 1 - Quasi-randomized *neural* networks (QRNs)

### QRNs applied to time series

- ▶ **Response**  $y$  = most recent time series observations
- ▶ **Covariates**  $X$  = time series lags
- ▶ **base learner**: can be any Machine Learning model
- ▶ **Multivariate forecasting case**: the base learner **shared** by all the time series

## 1 - 2 Uncertainty quantification in forecasting

### Point forecasts/Uncertainty quantification

- ▶ **Point forecasts:** cool, but not very informative. How wrong can we be, based on the assumptions that we made: **answers** how **“certain”** can we be about the forecast?
- ▶ **Uncertainty quantification** needed: **prediction intervals** and/or **predictive simulations**.
- ▶ **prediction intervals:** point forecast  $\pm$  a term (with a level of confidence)
- ▶ **predictive simulations:** future scenarios for the variables of interest



## 1 - 2 Uncertainty quantification in forecasting

### In `nnetsauce`

- ▶ Based on:
  - ▶ **Bayesian priors**
  - ▶ **In-sample residuals** = model fit - true observation on the *whole training set*
  - ▶ **Calibrated residuals** = model fit - true observation on a *held-out calibration set*
  - ▶ **(Vine) Copulas** (since `nnetsauce` v0.23.0)
- ▶ **Calibrated residuals** used in `nnetsauce` for methods based on **sequential split conformal prediction** (more on this later)

# 1 - 2 Uncertainty quantification in forecasting

Short focus on **sequential split conformal prediction** (see also <https://github.com/thierrymoudiki/2024-07-17-scp-block-bootstrap>)

## Sequential split conformal prediction + Block bootstrap

Data set:

AirPassengers

Transformation:

None

Forecasting method:

dynrm

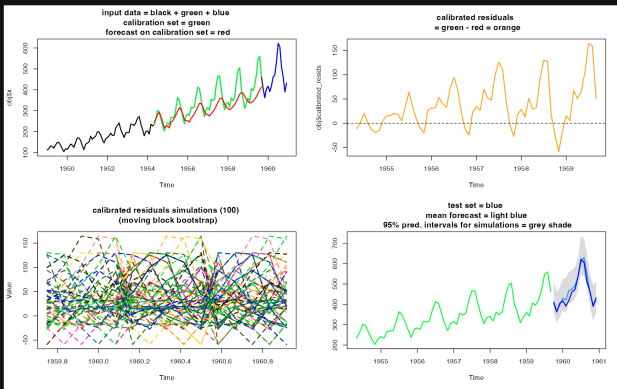
Block size:

5

### How it works:

- Split data into training, calibration and test set
- Point forecast on calibration set
- Obtain calibrated residuals
- Simulate calibrated residuals
- Point forecast on test set
- Calibrated residuals simulations + point forecast on test set

Plot of data:




# 1 - 2 Uncertainty quantification in forecasting

Short focus on *copulas* (source [vine-copula.org](http://vine-copula.org))

## Gumbel copula

(left  $\tau = .25$ , right:  $\tau = .75$ )

upper tail dependent



## Clayton copula

(left  $\tau = .25$ , right:  $\tau = .75$ )

lower tail dependent




Figure 2: Copulas

## 1 - 2 Uncertainty quantification in forecasting

### Recap

- ▶ In `nnetsauce` version 0.23.0:
  - ▶ Via a **Bayesian** base learner
  - ▶ Via a **conformalized** base learner
  - ▶ Via **in-sample** residuals for methods based on:
    - ▶ parametric residuals distribution inference (`gaussian`)
    - ▶ density estimation and simulation of residuals (`kde`)
    - ▶ bootstrap resampling (`bootstrap` and `block-bootstrap`)
    - ▶ vine copulas (`vine-*`)
  - ▶ Via **calibrated** residuals for methods based on **sequential split conformal prediction (SCP)** (`scp*-kde`, `scp*-bootstrap`, `scp*-block-bootstrap`, `scp*-vine-*`)

# Plan

- ▶ 1 - Key components of `nnetsauce` forecasting
  - ▶ 1 - 1 Quasi-randomized *neural* networks (QRNs)
  - ▶ 1 - 2 Uncertainty quantification in forecasting
- ▶ **2 - QRN forecasting with `nnetsauce`**
  - ▶ 2 - 1 `nnetsauce`'s description (Python version)
  - ▶ 2 - 2 Install+import Python packages (including `nnetsauce`)
  - ▶ 2 - 3 Import data for the demo
  - ▶ 2 - 4 Using the `fit + predict` interface
  - ▶ 2 - 5 Using GPUs
  - ▶ 2 - 6 Time series cross-validation
  - ▶ 2 - 7 AutoML with LazyMTS

## 2 - QRN forecasting with nnetsauce

### 2 - 1 nnetsauce's description (Python version)

- ▶ General-purpose Machine Learning using **Randomized** and **Quasi-Randomized *neural* networks**
  - ▶ GitHub: <https://github.com/Techtonique/nnetsauce>
  - ▶ PyPI: <https://pypi.org/project/nnetsauce/>
  - ▶ Conda: <https://anaconda.org/conda-forge/nnetsauce>
- ▶ **Tasks:**
  - ▶ Classification
  - ▶ Regression
  - ▶ **Univariate/Multivariate time series forecasting**

## 2 - 1 nnettsauce's description (Python version) (cont'd)

- ▶ **Simple interface** for each model:
  - ▶ `fit`: fitting model to training data
  - ▶ `predict`: model inference on unseen data
- ▶ **GPU** version optimizes matrices multiplication using **JAX** (not *magical*)
- ▶ Classes MTS and DeepMTS for **time series forecasting**
- ▶ DeepMTS seems to be more suited for *nearly* stationary data (but I encourage you to try and tell me)
- ▶ Automated Machine Learning (**AutoML**) with classes LazyMTS and LazyDeepMTS
- ▶ **Cross-validation**

## 2 - 2 Install+import Python packages (including nnetsauce)

```
pip install nnetsauce
```

```
pip install git+https://github.com/Techtonique/mlsauce.git
```

```
import nnetsauce as ns # import the package
import mlsauce as ms
import numpy as np
import pandas as pd
import seaborn as sns

from sklearn.linear_model import Ridge
from statsmodels.tsa.seasonal import STL

sns.set_theme(style="darkgrid")
```



## 2 - 3 Import data for the demo

### Input format (univariate time series)

date	value
1949-01-01	112.0
1949-02-01	118.0
1949-03-01	132.0
1949-04-01	129.0
1949-05-01	121.0
1949-06-01	135.0

## 2 - 3 Import data for the demo

### Input format (multivariate time series)


<b>date</b>	<b>Consumption</b>	<b>Income</b>
1970-01-01	0.615986218	0.972261043
1970-04-01	0.46037569	1.169084717
1970-07-01	0.876791423	1.55327055
1970-10-01	-0.274245141	-0.255272381
1971-01-01	1.897370758	1.987153628
1971-04-01	0.911992909	1.447334175

## 2 - 3 Import data for the demo

**Univariate: Monthly anti-diabetic drug sales in Australia from 1992 to 2008**

```
url = "https://raw.githubusercontent.com/Techtonique/"  
url += "datasets/main/time_series/univariate/"  
url += "a10.csv"  
df_a10 = pd.read_csv(url)  
df_a10.index = pd.DatetimeIndex(df_a10.date) # must have  
df_a10.drop(columns=['date'], inplace=True)
```

```
df_a10.plot()
```




## 2 - 3 Import data for the demo (cont'd)

### Multivariate: Heater vs Ice cream sales data set

```
url = "https://raw.githubusercontent.com/Techtonique/"
url += "datasets/main/time_series/multivariate/"
url += "ice_cream_vs_heater.csv"
df_temp = pd.read_csv(url)
df_temp.index = pd.DatetimeIndex(df_temp.date) # must have
# first other difference
df_icecream = df_temp.drop(columns=['date']).diff().\
dropna()
```

```
df_icecream.plot()
```



## 2 - 4 Using the `fit + predict` interface

**A few examples of probabilistic forecasting with `nnetsauce`:**

- ▶ Gaussian
- ▶ Bayesian (Gaussian prior on base learner)
- ▶ Kernel Density Estimation (KDE) and sequential split conformal prediction (SCP)
- ▶ Conformalized base learner: `TweedieRegressor` + SCP
- ▶ Vine Copula (combined with SCP)

Also use the **docs for exact spec.:**

<https://techtonique.github.io/nnetsauce/nnetsauce.html#MTS>

## 2 - 4 Using the fit + predict interface

### ► Gaussian

```
from sklearn.ensemble import BaggingRegressor


regr = ns.MTS(obj=BaggingRegressor(), # base learner
               type_pi="gaussian", # type of pred. int.
               lags=20, # number of time series lags
               show_progress=False)
regr.fit(df_icecream); # fit the model

regr.predict(h=30); # 30-steps ahead forecast
```



## 2 - 4 Using the fit + predict interface

```
regr.plot("heater", type_plot="pi") # plot pred. int.
```



## 2 - 4 Using the fit + predict interface

### ► Bayesian (Gaussian prior)


```
from sklearn.linear_model import BayesianRidge

regr = ns.MTS(obj=BayesianRidge(), # base learner
              lags=15, # no. of time series lags
              show_progress=False)
regr.fit(df_a10); # fit the model

regr.predict(h=40, return_std=True); # 40-steps ahead
```

## 2 - 4 Using the fit + predict interface

```
regr.plot(type_plot="pi")
```



## 2 - 4 Using the fit + predict interface

### ► SCP-KDE

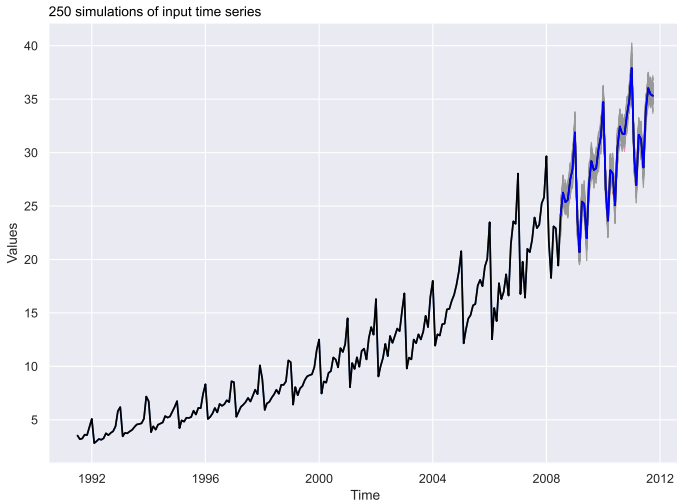
```
from sklearn.linear_model import LassoLarsIC

regr = ns.MTS(obj=LassoLarsIC(), # base learner
              type_pi="scp-kde", # type of pred. int.
              replications=250, # no. of sample paths
              kernel='gaussian', # density kernel
              lags=15, # no. of time series lags
              show_progress=False)
regr.fit(df_a10); # fit the model

regr.predict(h=40); # 40-steps ahead
```

## 2 - 4 Using the fit + predict interface

```
regr.plot(type_plot="spaghetti")
```



## 2 - 4 Using the fit + predict interface

### ► Conformalized base learner


```
from sklearn.linear_model import TweedieRegressor

obj0 = ns.PredictionInterval(obj=TweedieRegressor(),
                             method="splitconformal",
                             type_split="sequential",
                             level=95)

regr = ns.MTS(obj=obj0,
              lags=20,
              show_progress=False)
regr.fit(df_icecream);
regr.predict(h=30, return_pi=True);
```

## 2 - 4 Using the fit + predict interface

```
regr.plot("heater", type_plot="pi") # plot one series
```



## 2 - 4 Using the fit + predict interface

### ► Vine copula + sequential split conformal


```
import nnetsauce as ns
from sklearn.linear_model import TweedieRegressor

regr = ns.MTS(obj=TweedieRegressor(),
               lags=25, # no. of time series lags
               type_pi="scp-vine-tll", # vine copula spec.
               replications=250, # no. of sample paths
               show_progress=False)
regr.fit(df_icecream);
regr.predict(h=30);
```




## 2 - 4 Using the fit + predict interface

```
regr.plot("heater", type_plot="spaghetti")
```



## 2 - 4 Using the fit + predict interface

```
regr.plot("icecream", type_plot="spaghetti")
```



## 2 - 5 Using GPUs

- ▶ Public notebook on GitHub (<https://bit.ly/45RchgD>)
- ▶ Simulated multivariate time series: 100 series, 10000 observations
- ▶ Ran on **Kaggle notebooks**, with **accelerator GPU P100**

```
import numpy as np
import pandas as pd
import nnetsauce as ns
from sklearn.linear_model import Ridge
from time import time
```

```
df = generate_synthetic_mts(n_steps=10000,
                           n_series=100,
                           amplitude=40, seed=14531)
df_ = df.diff().dropna()
```

## 2 - 5 Using GPUs (cont'd)

- ▶ **Example 1** on CPU, using Ridge regression

```
regr = Ridge()
obj_MTS = ns.MTS(regr,
                 lags = 15,
                 n_hidden_features=5,
                 nodes_sim="uniform",
                 backend="cpu", # specify backend
                 verbose = 1)

start = time()
obj_MTS.fit(df_)
print(f"Elapsed: {time()-start}")
```

Elapsed: 64.46652388572693

## 2 - 5 Using GPUs (cont'd)

- **Example 2** on GPU (uses **JAX** behind the scenes), using Ridge regression

```
regr = Ridge()
obj_MTS = ns.MTS(regr,
                 lags = 15,
                 n_hidden_features=5,
                 nodes_sim="uniform",
                 backend="gpu", # specify backend
                 verbose = 1)

start = time()
obj_MTS.fit(df_)
print(f"Elapsed: {time()-start}")
```

Elapsed: 40.53069853782654

- **37% time gain**

## 2 - 5 Using GPUs (cont'd)

- ▶ **Example 3** on GPU (uses **JAX** behind the scenes), using Ridge regression on GPU (mlsauce implementation) **too**

```
regr = ms.RidgeRegressor(reg_lambda=1.0, backend="gpu")
obj_MTS = ns.MTS(regr,
                  lags = 15,
                  n_hidden_features=5,
                  nodes_sim="uniform",
                  backend="gpu", # specify backend
                  verbose = 1)

start = time()
obj_MTS.fit(df_)
print(f"Elapsed: {time()-start}")
```


Elapsed: 23.551459312438965

- ▶ **63% time gain**
- ▶ Can also use xgboost with tree\_method='gpu\_hist' e.g

## 2 - 6 Time series cross-validation


2 methods: **fixed window** and **increasing window**

**fixed window**



## 2 - 6 Time series cross-validation (cont'd)

### increasing window





## 2 - 6 Time series cross-validation (cont'd)

### Example in nnetsauce

```
import nnetsauce as ns
import numpy as np
import pandas as pd
import statsmodels.api as sm
from sklearn.linear_model import Ridge
from statsmodels.tsa.base.datetools import dates_from_str

# some example data
mdata = sm.datasets.macrodta.load_pandas().data
# prepare the dates index
dates = mdata[['year', 'quarter']].astype(int).astype(str)
quarterly = dates["year"] + "Q" + dates["quarter"]
quarterly = dates_from_str(quarterly)
mdata = mdata[['realgvt', 'tbilrate', 'cpi']]
mdata.index = pd.DatetimeIndex(quarterly)
data = np.log(mdata).diff().dropna()
```

## 2 - 6 Time series cross-validation (cont'd)

### Example in nnetnsauce (cont'd)

```
obj_MTS = ns.MTS(Ridge(), lags = 3,  
                  n_hidden_features=7,  
                  replications=100,  
                  seed=24, verbose = 0,  
                  type_pi="scp2-block-bootstrap",  
                  show_progress=False)  
cv = obj_MTS.cross_val_score(data,  
                              verbose = 0,  
                              initial_window=100,  
                              horizon=5,  
                              level=95,  
                              fixed_window=False, # True for rolling  
                              show_progress=False,  
                              scoring="coverage")[1]
```

## 2 - 6 Time series cross-validation (cont'd)

### **Example in nnetsauce (cont'd)**

Average coverage on 98 samples

```
print(cv.mean)
```

```
## 92.51700680272108
```

## 2 - 6 Time series cross-validation (cont'd)

### Example in nnetnsauce (cont'd)

```
obj_MTS = ns.MTS(Ridge(), lags = 3,
                  n_hidden_features=7,
                  replications=100,
                  seed=24, verbose = 0,
                  type_pi="scp2-block-bootstrap",
                  show_progress=False)
cv = obj_MTS.cross_val_score(data,
                              verbose = 0,
                              initial_window=100,
                              horizon=5,
                              level=95,
                              fixed_window=True, # False for increasing
                              show_progress=False,
                              scoring="coverage")[1]
```

## 2 - 6 Time series cross-validation (cont'd)

### **Example in nnetsauce (cont'd)**

Average coverage on 98 samples

```
print(cv.mean)
```

```
## 92.44897959183673
```

## 2 - 7 Automated Machine Learning (AutoML) with LazyMTS

```
# split data into training/testing set
n = data.shape[0]
max_idx_train = np.floor(n*0.9)
training_index = np.arange(0, max_idx_train)
testing_index = np.arange(max_idx_train, n)
df_train = data.iloc[training_index,:]
df_test = data.iloc[testing_index,:]
# Train + predict on 3 ML models
regr_mts = ns.LazyMTS(lags = 25,
                      type_pi="scp2-kde",
                      kernel="gaussian",
                      replications=250,
                      estimators=["Ridge",
                                  "ElasticNet",
                                  "RandomForestRegressor"],
                      show_progress=False);
models, predictions = regr_mts.fit(df_train, df_test);
```

## 2 - 7 Automated Machine Learning (AutoML) with LazyMTS (cont'd)

```
print(models[['RMSE', 'COVERAGE']])
```

##	RMSE	COVERAGE
## Model		
## MTS(ElasticNet)	0.20	87.30
## MTS(RandomForestRegressor)	0.20	88.89
## MTS(Ridge)	0.28	80.95

## 2 - 7 Automated Machine Learning (AutoML) with LazyMTS (cont'd)

```
regr_mts = ns.LazyMTS(lags = 25,  
                      type_pi="scp2-kde",  
                      kernel="gaussian",  
                      replications=250,  
                      estimators="all", # the change  
                      show_progress=False);  
models, predictions = regr_mts.fit(df_train, df_test);
```

```
##      0%|          | 0/32 [00:00<?, ?it/s]  3%|3          | 1
```



## 2 - 7 Automated Machine Learning (AutoML) with LazyMTS (cont'd)

```
print(models.sort_values(by='COVERAGE',  
ascending=False)[['COVERAGE']].head(5))
```

##	COVERAGE
## Model	
## MTS(AdaBoostRegressor)	92.06
## MTS(BaggingRegressor)	92.06
## MTS(ExtraTreeRegressor)	90.48
## MTS(KNeighborsRegressor)	90.48
## MTS(ElasticNetCV)	88.89